

LASSO Prediction on Churn Data

```
library(data.table) # Open the data.table library for use in this script
library(ggplot2)
library(glmnet)

set.seed(2093) # Set the seed so that all of our random operations produce the same results
```

CRM

CRM stands for Customer Retail Management. This refers to when the marketer has enough data to tailor the marketing experience of each customer. Here we will explore models that can be used to predict whether a customer will churn.

Data

The data is a development sample that consists of 100,000 randomly selected customers from the data base of a credit card company in Singapore. The random sample represents only a small percentage of the whole data base.

Customers are identified by a unique `customer_id`. The data also include demographic information on each customer, including whether their gender is male (`genderM`), their employment classification (`employXX`), income (`hhIncome`), household size (`hhSize`), whether they own a home (`homeOwner`), marital status (`married`). Additionally, some partial information on each customer's purchase history is included: the share of past months they have made late payments (`avgLatePayment`), the number of months they have been a customer of the credit card company (`nMonthsCust`), their average monthly bill (`avgMoBill`) and their credit score (`creditScore`).

Load the `churnCC` data set. We will use the `fread` command from the `data.table` package which opens text files like ".csv" much, much faster than base-R commands like `read.csv`.

```
churn_DT = fread("~/Dropbox/LargeRepoFiles/teaching/Teaching Data/output/Churn CC/data/churnCC.csv")
churn_DT[, .N] # .N is a special character for calling the number of observations in the data.table
churn_DT = churn_DT[runif(.N) < .01] # This creates a subset of the original data for testing/development
churn_DT[, .N]
```

Data inspection

Summarize some key aspects of the `churned` variable. In particular, what is the churn incidence, and the average churn value? Note that `data.table` offers a special character that may be used in the `j` position, `.N` which we will use here. `.N` reports the number of observations in the data set. When `.N` is combined with a `by` variable, then `data.table` reports the number of observations for each `by` group.

```
# Number of observations for each value of `churned`  
churn_DT[, .N, churned]
```

```
   churned    N  
1:      0 7650  
2:      1 2372
```

```
# Share of observations for which `churned` is zero or one.  
churn_DT[, .N / churn_DT[,.N], churned]
```

```
   churned      V1  
1:      0 0.7633207  
2:      1 0.2366793
```

```
# Average value of `churned` which matches the output above.-  
churn_DT[, mean(churned)]
```

```
[1] 0.2366793
```

Summarize the data

```
summary(churn_DT)
```

```

customerID      genderM      employSelfEmployed  employCrafts
Min.   :10000054  Min.   :0.0000  Min.   :0.00000  Min.   :0.0000
1st Qu.:10243584  1st Qu.:0.0000  1st Qu.:0.00000  1st Qu.:0.0000
Median :10499934  Median :1.0000  Median :0.00000  Median :0.0000
Mean   :10497001  Mean   :0.5019  Mean   :0.06735  Mean   :0.1107
3rd Qu.:10744425  3rd Qu.:1.0000  3rd Qu.:0.00000  3rd Qu.:0.0000
Max.   :10999971  Max.   :1.0000  Max.   :1.00000  Max.   :1.0000
employProfessional employClerical  employRetired  employStudent
Min.   :0.0000  Min.   :0.00000  Min.   :0.00000  Min.   :0.00000
1st Qu.:0.0000  1st Qu.:0.00000  1st Qu.:0.00000  1st Qu.:0.00000
Median :1.0000  Median :0.00000  Median :0.00000  Median :0.00000
Mean   :0.6488  Mean   :0.07633  Mean   :0.05448  Mean   :0.03233
3rd Qu.:1.0000  3rd Qu.:0.00000  3rd Qu.:0.00000  3rd Qu.:0.00000
Max.   :1.0000  Max.   :1.00000  Max.   :1.00000  Max.   :1.00000

hhIncome      hhSize      homeOwner      married
Min.   :      0  Min.   :1.000  Min.   :0.0000  Min.   :0.0000
1st Qu.: 40000  1st Qu.:5.000  1st Qu.:1.0000  1st Qu.:0.0000
Median : 60000  Median :5.000  Median :1.0000  Median :1.0000
Mean   : 60010  Mean   :4.561  Mean   :0.8639  Mean   :0.7094
3rd Qu.: 80000  3rd Qu.:5.000  3rd Qu.:1.0000  3rd Qu.:1.0000
Max.   :180000  Max.   :5.000  Max.   :1.0000  Max.   :1.0000

avgLatePayment  nMonthsCust      avgMoBill      creditScore
Min.   :3.901e-05  Min.   : 0.000  Min.   : 10.0  Min.   :200.0
1st Qu.:5.016e-02  1st Qu.: 3.000  1st Qu.: 80.0  1st Qu.:440.0
Median :1.008e-01  Median : 4.000  Median :100.0  Median :520.0
Mean   :1.003e-01  Mean   : 3.977  Mean   :100.2  Mean   :520.2
3rd Qu.:1.506e-01  3rd Qu.: 5.000  3rd Qu.:120.0  3rd Qu.:600.0
Max.   :2.000e-01  Max.   :14.000  Max.   :240.0  Max.   :800.0

churned
Min.   :0.0000
1st Qu.:0.0000
Median :0.0000
Mean   :0.2367
3rd Qu.:0.0000
Max.   :1.0000

```

In this analysis we would like to predict who will churn, or the **churned** variable. We should take a minute to think about which of these variables we want to use to predict **churned**. Most of the variables appear to have the potential to predict **churned**, however, there are a few that should not be included as predictive variables.

First, **customerID** should have no relationship to whether a customer churned; these values are just a way to track customers. Additionally, even if **customerID** was somehow able to predict **churned** we would not want to use it because when we apply our model to predict churn for a different group of customers, their **customerID** values will be different. In this case, a model that includes **customerID** will not be helpful. In general, we want to restrict our predictor variables to those that will be available to us, with the same types of values.

Add additional variables that are random numbers

To help illustrate the value of LASSO regression, we are going to add several new variables that are random and do not help predict churned. (Normally, we would not do this.)

```
# This syntax is more advanced data.table and not a core part of this course  
# This code creates many "garbage variables" that are just random noise.  
  
garbNum = 300  
garbVars = paste0("X", 1:garbNum)  
churn_DT[, (garbVars) := 1]  
churn_DT[, (garbVars) := lapply(.SD, function (x) runif(.N) ), .SDcols = garbVars]
```

Holdout sample

Because we want to make predictions and assess the quality of those predictions, we need to produce a **confusion matrix** for our predictions. To produce this matrix we need observations which include both our model's predictions and the actual outcomes. How do we generate such data? Well, we already have observations with actual outcomes, so we just need to add model predictions.

However, we have to be careful not to use the same data for estimating the model and for producing the **confusion matrix**. Using the same data leads to **overfit** a topic we will discuss further. For now, suffice it to say that we want to evaluate the predictions our model on data we have never seen before. After all, this is how we plan to use the predictions, so this is how we should evaluate their quality.

A good way to estimate the **confusion matrix** is by generating a **holdout sample**. A holdout sample is a random subset of the data we have. We will “estimate” or “train” our model using the “training data.” Then, with our model estimated, we will test its predictions using the holdout data. The holdout data will be used to produce the confusion matrix.

```
churn_DT[, holdout := sample(0:1, size=.N, prob=c(0.5, 0.5), replace=TRUE)]
```

Estimate the model

For now we will use a standard logit model to make our predictions. This step is also known as *training* the model or *calibrating* the model. In later classes we will introduce other methods for making predictions. First we create a vector of the names of variables we don't want included in the regression. Placing them in `noRegVar` allows us to refer to these variables later.

```
# Variables that are not to be used for estimation  
names(churn_DT)
```

[1] "customerID"	"genderM"	"employSelfEmployed"
[4] "employCrafts"	"employProfessional"	"employClerical"
[7] "employRetired"	"employStudent"	"hhIncome"
[10] "hhSize"	"homeOwner"	"married"
[13] "avgLatePayment"	"nMonthsCust"	"avgMoBill"
[16] "creditScore"	"churned"	"X1"
[19] "X2"	"X3"	"X4"
[22] "X5"	"X6"	"X7"
[25] "X8"	"X9"	"X10"
[28] "X11"	"X12"	"X13"
[31] "X14"	"X15"	"X16"
[34] "X17"	"X18"	"X19"
[37] "X20"	"X21"	"X22"
[40] "X23"	"X24"	"X25"
[43] "X26"	"X27"	"X28"

[46]	"X29"	"X30"	"X31"
[49]	"X32"	"X33"	"X34"
[52]	"X35"	"X36"	"X37"
[55]	"X38"	"X39"	"X40"
[58]	"X41"	"X42"	"X43"
[61]	"X44"	"X45"	"X46"
[64]	"X47"	"X48"	"X49"
[67]	"X50"	"X51"	"X52"
[70]	"X53"	"X54"	"X55"
[73]	"X56"	"X57"	"X58"
[76]	"X59"	"X60"	"X61"
[79]	"X62"	"X63"	"X64"
[82]	"X65"	"X66"	"X67"
[85]	"X68"	"X69"	"X70"
[88]	"X71"	"X72"	"X73"
[91]	"X74"	"X75"	"X76"
[94]	"X77"	"X78"	"X79"
[97]	"X80"	"X81"	"X82"
[100]	"X83"	"X84"	"X85"
[103]	"X86"	"X87"	"X88"
[106]	"X89"	"X90"	"X91"
[109]	"X92"	"X93"	"X94"
[112]	"X95"	"X96"	"X97"
[115]	"X98"	"X99"	"X100"
[118]	"X101"	"X102"	"X103"
[121]	"X104"	"X105"	"X106"
[124]	"X107"	"X108"	"X109"
[127]	"X110"	"X111"	"X112"
[130]	"X113"	"X114"	"X115"
[133]	"X116"	"X117"	"X118"
[136]	"X119"	"X120"	"X121"
[139]	"X122"	"X123"	"X124"
[142]	"X125"	"X126"	"X127"
[145]	"X128"	"X129"	"X130"
[148]	"X131"	"X132"	"X133"
[151]	"X134"	"X135"	"X136"
[154]	"X137"	"X138"	"X139"
[157]	"X140"	"X141"	"X142"
[160]	"X143"	"X144"	"X145"
[163]	"X146"	"X147"	"X148"
[166]	"X149"	"X150"	"X151"
[169]	"X152"	"X153"	"X154"
[172]	"X155"	"X156"	"X157"
[175]	"X158"	"X159"	"X160"
[178]	"X161"	"X162"	"X163"
[181]	"X164"	"X165"	"X166"
[184]	"X167"	"X168"	"X169"
[187]	"X170"	"X171"	"X172"
[190]	"X173"	"X174"	"X175"
[193]	"X176"	"X177"	"X178"
[196]	"X179"	"X180"	"X181"
[199]	"X182"	"X183"	"X184"
[202]	"X185"	"X186"	"X187"
[205]	"X188"	"X189"	"X190"

[208]	"X191"	"X192"	"X193"
[211]	"X194"	"X195"	"X196"
[214]	"X197"	"X198"	"X199"
[217]	"X200"	"X201"	"X202"
[220]	"X203"	"X204"	"X205"
[223]	"X206"	"X207"	"X208"
[226]	"X209"	"X210"	"X211"
[229]	"X212"	"X213"	"X214"
[232]	"X215"	"X216"	"X217"
[235]	"X218"	"X219"	"X220"
[238]	"X221"	"X222"	"X223"
[241]	"X224"	"X225"	"X226"
[244]	"X227"	"X228"	"X229"
[247]	"X230"	"X231"	"X232"
[250]	"X233"	"X234"	"X235"
[253]	"X236"	"X237"	"X238"
[256]	"X239"	"X240"	"X241"
[259]	"X242"	"X243"	"X244"
[262]	"X245"	"X246"	"X247"
[265]	"X248"	"X249"	"X250"
[268]	"X251"	"X252"	"X253"
[271]	"X254"	"X255"	"X256"
[274]	"X257"	"X258"	"X259"
[277]	"X260"	"X261"	"X262"
[280]	"X263"	"X264"	"X265"
[283]	"X266"	"X267"	"X268"
[286]	"X269"	"X270"	"X271"
[289]	"X272"	"X273"	"X274"
[292]	"X275"	"X276"	"X277"
[295]	"X278"	"X279"	"X280"
[298]	"X281"	"X282"	"X283"
[301]	"X284"	"X285"	"X286"
[304]	"X287"	"X288"	"X289"
[307]	"X290"	"X291"	"X292"
[310]	"X293"	"X294"	"X295"
[313]	"X296"	"X297"	"X298"
[316]	"X299"	"X300"	"holdout"

```
noRegVars = c("holdout", "customerID")

churn.fit.base = glm(churned ~ ., family="binomial",
  data=churn_DT[holdout==0, -c(..noRegVars, ..garbVars)] )

# Estimate a logit model with garbage variables
churn.fit.garb = glm(churned ~ ., family="binomial",
  data=churn_DT[holdout==0, -..noRegVars] )
```

We use the `glm` command to estimate a logit model with the setting `family="binomial"`. Note that the first entry in `gml` is the formula for the regression we want to estimate. `churned ~ .` indicates that we want to predict `churned` using all of the other variables in the data set, represented by `..`.

For the data set, we entered into to `glm` we have used some `data.table` commands to produce a subset of the original data. First, we selected only observations that are *not* in the holdout sample using `holdout==0`. In other words, we have selected the “training” data to estimate the model. Second, we removed the variables included in `noRegVars` using some `data.table` syntax. The minus sign `-` indicates that we are removing

the variables. The double dots, .. indicate that we are entering a vector of variable names and want to (de-)select those variables from the data set.

Inspect the estimates

```
summary(churn.fit.garb)
```

Call:

```
glm(formula = churned ~ ., family = "binomial", data = churn_DT[holdout ==  
0, -..noRegVars])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9476	-0.6948	-0.4356	-0.1656	3.1015

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-3.537e+00	1.261e+00	-2.805	0.00504	**
genderM	-1.752e-01	7.793e-02	-2.249	0.02454	*
employSelfEmployed	-3.130e-01	3.964e-01	-0.790	0.42969	
employCrafts	-4.044e-01	3.871e-01	-1.045	0.29620	
employProfessional	-6.853e-01	3.734e-01	-1.835	0.06644	.
employClerical	-3.302e-01	3.937e-01	-0.839	0.40170	
employRetired	-2.941e-02	4.010e-01	-0.073	0.94153	
employStudent	-5.278e-01	4.297e-01	-1.228	0.21929	
hhIncome	-4.135e-06	1.585e-06	-2.609	0.00907	**
hhSize	3.907e-02	3.749e-02	1.042	0.29736	
homeOwner	1.021e-01	1.164e-01	0.877	0.38051	
married	-1.388e-01	8.572e-02	-1.620	0.10534	
avgLatePayment	5.782e-01	6.689e-01	0.865	0.38731	
nMonthsCust	9.718e-03	1.949e-02	0.499	0.61811	
avgMoBill	3.242e-02	1.396e-03	23.223	< 2e-16	***
creditScore	-4.242e-04	3.516e-04	-1.207	0.22761	
X1	-2.765e-02	1.354e-01	-0.204	0.83819	
X2	-3.111e-01	1.352e-01	-2.301	0.02138	*
X3	6.729e-02	1.342e-01	0.501	0.61620	
X4	-1.848e-01	1.369e-01	-1.350	0.17692	
X5	-3.662e-03	1.342e-01	-0.027	0.97823	
X6	1.772e-01	1.347e-01	1.316	0.18834	
X7	1.272e-02	1.354e-01	0.094	0.92517	
X8	-1.852e-01	1.345e-01	-1.377	0.16856	
X9	-1.396e-01	1.368e-01	-1.020	0.30753	
X10	-5.693e-02	1.360e-01	-0.419	0.67544	
X11	-2.117e-01	1.362e-01	-1.555	0.11999	
X12	1.671e-02	1.344e-01	0.124	0.90105	
X13	-3.274e-02	1.334e-01	-0.245	0.80612	
X14	-1.335e-01	1.339e-01	-0.997	0.31867	
X15	2.141e-02	1.374e-01	0.156	0.87613	
X16	-2.615e-01	1.345e-01	-1.944	0.05188	.
X17	2.513e-02	1.371e-01	0.183	0.85459	
X18	2.054e-01	1.353e-01	1.518	0.12905	
X19	2.595e-01	1.354e-01	1.917	0.05525	.
X20	1.114e-01	1.362e-01	0.818	0.41337	
X21	-7.357e-02	1.352e-01	-0.544	0.58633	

X22	9.944e-02	1.353e-01	0.735	0.46225
X23	3.903e-02	1.327e-01	0.294	0.76870
X24	3.275e-01	1.347e-01	2.432	0.01502 *
X25	1.543e-01	1.339e-01	1.153	0.24906
X26	-1.220e-01	1.351e-01	-0.903	0.36649
X27	-7.936e-02	1.362e-01	-0.583	0.56015
X28	-1.167e-01	1.360e-01	-0.857	0.39119
X29	1.250e-01	1.359e-01	0.920	0.35747
X30	-8.931e-02	1.358e-01	-0.658	0.51062
X31	9.830e-02	1.345e-01	0.731	0.46490
X32	2.600e-02	1.357e-01	0.192	0.84807
X33	-4.366e-02	1.351e-01	-0.323	0.74663
X34	-1.921e-01	1.350e-01	-1.422	0.15488
X35	1.584e-01	1.329e-01	1.192	0.23308
X36	1.080e-01	1.348e-01	0.801	0.42304
X37	5.450e-02	1.342e-01	0.406	0.68458
X38	-4.427e-03	1.359e-01	-0.033	0.97401
X39	-4.531e-02	1.356e-01	-0.334	0.73826
X40	2.142e-01	1.364e-01	1.570	0.11632
X41	9.877e-02	1.350e-01	0.732	0.46431
X42	-2.809e-02	1.344e-01	-0.209	0.83441
X43	-1.619e-01	1.361e-01	-1.190	0.23423
X44	-6.595e-02	1.359e-01	-0.485	0.62754
X45	2.759e-03	1.352e-01	0.020	0.98372
X46	3.426e-01	1.382e-01	2.479	0.01318 *
X47	-1.789e-01	1.337e-01	-1.338	0.18084
X48	8.038e-02	1.353e-01	0.594	0.55232
X49	2.036e-01	1.337e-01	1.523	0.12783
X50	4.018e-03	1.366e-01	0.029	0.97653
X51	2.591e-02	1.369e-01	0.189	0.84992
X52	-9.330e-02	1.359e-01	-0.687	0.49239
X53	-1.223e-02	1.349e-01	-0.091	0.92776
X54	2.443e-02	1.347e-01	0.181	0.85608
X55	1.639e-01	1.355e-01	1.209	0.22660
X56	1.704e-01	1.352e-01	1.261	0.20734
X57	3.320e-02	1.349e-01	0.246	0.80560
X58	-1.738e-01	1.343e-01	-1.294	0.19577
X59	2.159e-01	1.335e-01	1.617	0.10595
X60	-1.317e-01	1.353e-01	-0.974	0.33018
X61	9.889e-02	1.336e-01	0.740	0.45900
X62	2.310e-02	1.360e-01	0.170	0.86515
X63	1.947e-01	1.346e-01	1.446	0.14817
X64	5.929e-02	1.359e-01	0.436	0.66263
X65	-1.779e-01	1.340e-01	-1.328	0.18428
X66	4.269e-02	1.375e-01	0.310	0.75622
X67	-6.873e-02	1.348e-01	-0.510	0.61026
X68	-1.561e-02	1.363e-01	-0.115	0.90878
X69	-9.168e-02	1.344e-01	-0.682	0.49527
X70	3.430e-01	1.362e-01	2.519	0.01178 *
X71	-1.370e-01	1.362e-01	-1.006	0.31443
X72	8.402e-02	1.356e-01	0.620	0.53553
X73	-1.168e-01	1.365e-01	-0.856	0.39215
X74	8.782e-02	1.350e-01	0.651	0.51525
X75	6.272e-02	1.377e-01	0.455	0.64886

X76	7.533e-02	1.351e-01	0.558	0.57709
X77	2.167e-01	1.352e-01	1.602	0.10913
X78	-8.331e-02	1.331e-01	-0.626	0.53127
X79	7.415e-02	1.355e-01	0.547	0.58410
X80	1.445e-01	1.329e-01	1.087	0.27691
X81	1.443e-01	1.354e-01	1.066	0.28656
X82	2.105e-01	1.358e-01	1.550	0.12112
X83	1.741e-01	1.349e-01	1.290	0.19705
X84	2.701e-02	1.364e-01	0.198	0.84304
X85	2.875e-03	1.347e-01	0.021	0.98298
X86	-2.175e-02	1.343e-01	-0.162	0.87133
X87	-8.284e-02	1.380e-01	-0.600	0.54831
X88	3.203e-01	1.352e-01	2.369	0.01785 *
X89	-6.149e-02	1.350e-01	-0.455	0.64887
X90	9.841e-02	1.351e-01	0.728	0.46652
X91	-1.398e-01	1.355e-01	-1.032	0.30208
X92	4.724e-02	1.369e-01	0.345	0.73002
X93	-1.805e-01	1.327e-01	-1.360	0.17378
X94	-1.824e-01	1.342e-01	-1.359	0.17417
X95	2.789e-01	1.349e-01	2.068	0.03862 *
X96	2.724e-01	1.363e-01	1.999	0.04563 *
X97	-8.091e-02	1.361e-01	-0.595	0.55214
X98	-5.001e-02	1.369e-01	-0.365	0.71492
X99	-5.776e-02	1.383e-01	-0.418	0.67625
X100	-3.793e-03	1.351e-01	-0.028	0.97761
X101	1.979e-01	1.342e-01	1.474	0.14043
X102	2.936e-01	1.347e-01	2.179	0.02932 *
X103	-2.271e-01	1.354e-01	-1.677	0.09354 .
X104	1.109e-01	1.335e-01	0.831	0.40593
X105	5.727e-02	1.356e-01	0.422	0.67272
X106	9.075e-02	1.361e-01	0.667	0.50488
X107	4.682e-02	1.378e-01	0.340	0.73396
X108	1.555e-01	1.361e-01	1.143	0.25312
X109	-5.101e-02	1.380e-01	-0.370	0.71162
X110	1.184e-01	1.341e-01	0.883	0.37713
X111	-1.372e-01	1.336e-01	-1.027	0.30443
X112	-9.802e-02	1.361e-01	-0.720	0.47151
X113	9.058e-02	1.348e-01	0.672	0.50148
X114	5.747e-03	1.367e-01	0.042	0.96646
X115	3.999e-02	1.337e-01	0.299	0.76483
X116	5.331e-02	1.353e-01	0.394	0.69363
X117	9.449e-03	1.343e-01	0.070	0.94392
X118	7.377e-02	1.368e-01	0.539	0.58958
X119	2.029e-01	1.350e-01	1.503	0.13285
X120	2.152e-01	1.352e-01	1.592	0.11147
X121	-4.866e-02	1.346e-01	-0.361	0.71777
X122	-6.296e-02	1.349e-01	-0.467	0.64065
X123	-6.825e-02	1.339e-01	-0.510	0.61014
X124	-1.560e-02	1.374e-01	-0.113	0.90964
X125	-2.022e-01	1.368e-01	-1.479	0.13922
X126	1.231e-01	1.349e-01	0.913	0.36139
X127	-1.730e-01	1.346e-01	-1.285	0.19872
X128	-5.017e-02	1.352e-01	-0.371	0.71047
X129	7.600e-02	1.355e-01	0.561	0.57494

X130	1.302e-01	1.363e-01	0.955	0.33933
X131	-7.610e-02	1.350e-01	-0.564	0.57308
X132	2.339e-02	1.359e-01	0.172	0.86335
X133	-6.303e-02	1.343e-01	-0.469	0.63889
X134	-5.348e-05	1.357e-01	0.000	0.99969
X135	-6.914e-02	1.354e-01	-0.511	0.60957
X136	-3.773e-01	1.356e-01	-2.783	0.00538 **
X137	-1.381e-01	1.345e-01	-1.027	0.30436
X138	-1.549e-01	1.364e-01	-1.135	0.25625
X139	-2.884e-01	1.329e-01	-2.170	0.02998 *
X140	2.922e-02	1.328e-01	0.220	0.82590
X141	-2.116e-01	1.335e-01	-1.585	0.11300
X142	1.510e-01	1.337e-01	1.129	0.25874
X143	-2.394e-01	1.357e-01	-1.764	0.07766 .
X144	-2.581e-01	1.355e-01	-1.905	0.05676 .
X145	8.600e-02	1.371e-01	0.627	0.53060
X146	1.761e-01	1.342e-01	1.312	0.18955
X147	-1.284e-02	1.348e-01	-0.095	0.92409
X148	-6.447e-02	1.352e-01	-0.477	0.63332
X149	-1.832e-01	1.356e-01	-1.351	0.17681
X150	-1.070e-01	1.375e-01	-0.778	0.43644
X151	1.154e-01	1.362e-01	0.847	0.39700
X152	1.316e-01	1.350e-01	0.975	0.32956
X153	-1.630e-01	1.347e-01	-1.210	0.22630
X154	1.600e-01	1.359e-01	1.177	0.23914
X155	1.406e-01	1.348e-01	1.044	0.29669
X156	-4.950e-02	1.358e-01	-0.364	0.71549
X157	-1.349e-01	1.342e-01	-1.005	0.31468
X158	2.067e-02	1.353e-01	0.153	0.87852
X159	3.584e-01	1.393e-01	2.572	0.01010 *
X160	-8.826e-02	1.350e-01	-0.654	0.51324
X161	-3.438e-01	1.340e-01	-2.565	0.01032 *
X162	6.777e-02	1.357e-01	0.499	0.61746
X163	1.503e-01	1.374e-01	1.094	0.27395
X164	-9.397e-02	1.365e-01	-0.689	0.49105
X165	-6.023e-02	1.345e-01	-0.448	0.65428
X166	2.557e-02	1.369e-01	0.187	0.85190
X167	-2.447e-01	1.341e-01	-1.825	0.06800 .
X168	1.789e-01	1.352e-01	1.324	0.18556
X169	3.149e-01	1.344e-01	2.344	0.01909 *
X170	2.024e-01	1.356e-01	1.493	0.13540
X171	6.335e-02	1.360e-01	0.466	0.64144
X172	2.204e-02	1.360e-01	0.162	0.87125
X173	-2.009e-01	1.343e-01	-1.496	0.13467
X174	-1.721e-01	1.362e-01	-1.264	0.20632
X175	-8.245e-02	1.353e-01	-0.609	0.54230
X176	-2.279e-01	1.371e-01	-1.662	0.09642 .
X177	2.816e-02	1.368e-01	0.206	0.83687
X178	-1.184e-01	1.347e-01	-0.879	0.37920
X179	-1.546e-01	1.340e-01	-1.154	0.24852
X180	-4.854e-02	1.352e-01	-0.359	0.71961
X181	1.095e-01	1.346e-01	0.814	0.41593
X182	-5.384e-02	1.350e-01	-0.399	0.69010
X183	2.687e-01	1.354e-01	1.985	0.04717 *

X184	-2.407e-01	1.358e-01	-1.773	0.07629	.
X185	-2.421e-02	1.354e-01	-0.179	0.85810	
X186	-1.033e-01	1.343e-01	-0.769	0.44175	
X187	2.220e-02	1.369e-01	0.162	0.87112	
X188	-6.012e-02	1.367e-01	-0.440	0.66002	
X189	1.338e-01	1.356e-01	0.987	0.32355	
X190	-1.052e-01	1.365e-01	-0.771	0.44098	
X191	-2.343e-01	1.340e-01	-1.749	0.08028	.
X192	9.628e-02	1.359e-01	0.709	0.47860	
X193	-5.260e-02	1.355e-01	-0.388	0.69792	
X194	-1.919e-01	1.357e-01	-1.414	0.15736	
X195	-1.551e-01	1.356e-01	-1.144	0.25272	
X196	3.749e-01	1.339e-01	2.799	0.00513	**
X197	1.466e-01	1.351e-01	1.085	0.27784	
X198	-2.490e-02	1.347e-01	-0.185	0.85334	
X199	-4.442e-02	1.337e-01	-0.332	0.73977	
X200	1.096e-01	1.361e-01	0.805	0.42080	
X201	-8.679e-03	1.360e-01	-0.064	0.94913	
X202	2.938e-02	1.367e-01	0.215	0.82988	
X203	-1.396e-01	1.340e-01	-1.041	0.29770	
X204	4.186e-02	1.372e-01	0.305	0.76021	
X205	-2.550e-02	1.310e-01	-0.195	0.84566	
X206	8.049e-02	1.348e-01	0.597	0.55037	
X207	2.970e-01	1.367e-01	2.173	0.02980	*
X208	-4.290e-02	1.340e-01	-0.320	0.74892	
X209	4.961e-03	1.342e-01	0.037	0.97050	
X210	-1.150e-02	1.339e-01	-0.086	0.93156	
X211	7.622e-02	1.356e-01	0.562	0.57398	
X212	1.305e-01	1.358e-01	0.961	0.33664	
X213	-1.462e-01	1.361e-01	-1.075	0.28254	
X214	8.141e-02	1.351e-01	0.603	0.54664	
X215	1.021e-01	1.354e-01	0.754	0.45094	
X216	5.403e-02	1.351e-01	0.400	0.68915	
X217	1.047e-01	1.342e-01	0.781	0.43505	
X218	6.697e-02	1.384e-01	0.484	0.62857	
X219	-1.301e-01	1.358e-01	-0.958	0.33819	
X220	9.297e-02	1.342e-01	0.693	0.48836	
X221	1.533e-02	1.349e-01	0.114	0.90954	
X222	-5.824e-02	1.341e-01	-0.434	0.66412	
X223	-5.004e-02	1.365e-01	-0.367	0.71395	
X224	-2.098e-01	1.362e-01	-1.541	0.12339	
X225	1.886e-01	1.357e-01	1.390	0.16450	
X226	3.033e-02	1.350e-01	0.225	0.82222	
X227	-1.000e-01	1.338e-01	-0.747	0.45478	
X228	-2.412e-01	1.361e-01	-1.773	0.07626	.
X229	-1.776e-01	1.357e-01	-1.309	0.19062	
X230	2.791e-02	1.364e-01	0.205	0.83790	
X231	-1.416e-01	1.355e-01	-1.045	0.29623	
X232	-1.635e-01	1.353e-01	-1.209	0.22669	
X233	1.920e-03	1.355e-01	0.014	0.98870	
X234	-1.036e-01	1.356e-01	-0.764	0.44465	
X235	2.863e-02	1.331e-01	0.215	0.82968	
X236	-1.071e-01	1.349e-01	-0.793	0.42752	
X237	1.001e-01	1.375e-01	0.728	0.46660	

X238	-3.696e-01	1.369e-01	-2.700	0.00694	**
X239	-2.478e-01	1.347e-01	-1.840	0.06584	.
X240	-2.434e-01	1.348e-01	-1.806	0.07088	.
X241	2.603e-02	1.357e-01	0.192	0.84791	
X242	-1.484e-01	1.352e-01	-1.098	0.27235	
X243	1.799e-01	1.356e-01	1.327	0.18458	
X244	-1.907e-02	1.349e-01	-0.141	0.88763	
X245	1.120e-01	1.359e-01	0.824	0.40977	
X246	6.810e-02	1.361e-01	0.500	0.61675	
X247	-4.389e-01	1.375e-01	-3.193	0.00141	**
X248	2.010e-01	1.344e-01	1.495	0.13482	
X249	-2.590e-02	1.363e-01	-0.190	0.84934	
X250	-3.207e-01	1.335e-01	-2.403	0.01626	*
X251	-2.310e-02	1.360e-01	-0.170	0.86509	
X252	9.967e-03	1.360e-01	0.073	0.94158	
X253	-5.352e-02	1.343e-01	-0.398	0.69029	
X254	4.234e-01	1.350e-01	3.137	0.00171	**
X255	-2.090e-01	1.361e-01	-1.536	0.12462	
X256	-8.626e-02	1.351e-01	-0.638	0.52332	
X257	2.149e-02	1.328e-01	0.162	0.87144	
X258	1.929e-01	1.355e-01	1.424	0.15452	
X259	-8.838e-03	1.340e-01	-0.066	0.94742	
X260	-9.629e-02	1.348e-01	-0.714	0.47495	
X261	-2.989e-02	1.361e-01	-0.220	0.82613	
X262	-2.606e-01	1.356e-01	-1.922	0.05461	.
X263	-2.954e-01	1.340e-01	-2.204	0.02755	*
X264	1.966e-01	1.354e-01	1.452	0.14647	
X265	1.427e-01	1.363e-01	1.046	0.29534	
X266	5.816e-02	1.359e-01	0.428	0.66858	
X267	-2.174e-01	1.355e-01	-1.604	0.10866	
X268	6.107e-03	1.335e-01	0.046	0.96352	
X269	-6.670e-02	1.348e-01	-0.495	0.62067	
X270	3.654e-02	1.371e-01	0.267	0.78974	
X271	-1.101e-01	1.343e-01	-0.820	0.41229	
X272	4.710e-02	1.359e-01	0.347	0.72889	
X273	-4.377e-02	1.362e-01	-0.321	0.74786	
X274	3.096e-01	1.345e-01	2.301	0.02138	*
X275	-1.129e-01	1.361e-01	-0.830	0.40650	
X276	1.841e-01	1.363e-01	1.350	0.17692	
X277	-3.917e-02	1.352e-01	-0.290	0.77202	
X278	-9.402e-02	1.340e-01	-0.702	0.48275	
X279	-2.305e-01	1.355e-01	-1.701	0.08896	.
X280	6.473e-02	1.363e-01	0.475	0.63480	
X281	6.566e-02	1.350e-01	0.486	0.62663	
X282	-7.138e-02	1.353e-01	-0.528	0.59784	
X283	1.968e-01	1.354e-01	1.454	0.14603	
X284	2.173e-01	1.348e-01	1.612	0.10687	
X285	3.216e-02	1.379e-01	0.233	0.81556	
X286	2.506e-02	1.343e-01	0.187	0.85203	
X287	-4.410e-02	1.345e-01	-0.328	0.74294	
X288	1.182e-01	1.371e-01	0.862	0.38850	
X289	-1.698e-01	1.340e-01	-1.267	0.20503	
X290	-2.758e-01	1.356e-01	-2.035	0.04189	*
X291	-3.469e-02	1.358e-01	-0.255	0.79834	

```

X292          1.148e-01  1.351e-01   0.849  0.39561
X293          1.044e-01  1.338e-01   0.780  0.43533
X294          2.119e-02  1.352e-01   0.157  0.87549
X295         -4.318e-01  1.349e-01  -3.202  0.00137 **
X296         -2.274e-02  1.332e-01  -0.171  0.86444
X297          7.367e-02  1.356e-01   0.543  0.58683
X298         -1.379e-01  1.360e-01  -1.014  0.31049
X299          2.090e-01  1.354e-01   1.544  0.12256
X300          8.687e-02  1.344e-01   0.646  0.51814
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 5520.5  on 5021  degrees of freedom
Residual deviance: 4431.0  on 4706  degrees of freedom
AIC: 5063

```

Number of Fisher Scoring iterations: 5

Based on the coefficient estimates and their z-values, only a few of the variables we used in the model are statistically significant. This is okay, we didn't have clear sense of which variables would be important. Additionally, because we are primarily interested in making prediction with this model, we are not especially interested in the coefficient estimates.

Make predictions in a Holdout Sample

Now that we have estimated the model, we can generate the confusion matrix. Let's create a new data.table just for this purpose:

```

churn_HO = churn_DT[holdout==1]
chHOMat = model.matrix( churned ~ ., churn_HO[, -..noRegVars])

```

Next, let's make predictions for this data set. Recall that none of this data was used to estimate the model, which is important for the validity of the predictions. We want to test our "out-of-sample" predictions (predictions on data other than the training data) since that is how we will use the model.

```

churn_HO[, predLogitProbBase := predict(churn.fit.base, newdata=churn_HO, type = "response")]
churn_HO[, predLogitProbGarb := predict(churn.fit.garb, newdata=churn_HO, type = "response")]

```


This step used the estimated logit model to predict the probability that each customer in the holdout data would churn. Let's have a look at the distribution of these probabilities

```

churn_HO[, hist(predLogitProbGarb)]

```

Histogram of predLogitProbGarb



\$breaks

```
[1] 0.00 0.05 0.10 0.15 0.20 0.25 0.30 0.35 0.40 0.45 0.50 0.55 0.60 0.65 0.70  
[16] 0.75 0.80 0.85 0.90 0.95 1.00
```

\$counts

```
[1] 668 872 691 533 402 359 257 245 186 143 136 110 102 76 77 56 37 28 16  
[20] 6
```

\$density

```
[1] 2.672 3.488 2.764 2.132 1.608 1.436 1.028 0.980 0.744 0.572 0.544 0.440  
[13] 0.408 0.304 0.308 0.224 0.148 0.112 0.064 0.024
```

\$mids

```
[1] 0.025 0.075 0.125 0.175 0.225 0.275 0.325 0.375 0.425 0.475 0.525 0.575  
[13] 0.625 0.675 0.725 0.775 0.825 0.875 0.925 0.975
```

\$xname

```
[1] "predLogitProbGarb"
```

\$equidist

```
[1] TRUE
```

attr("class")

```
[1] "histogram"
```

Plot the Profit Curve

A function for calculating expected profits.

The following function, called `expProfit` provides a way to compute the expected profits associated with targeting different portions of the population. It should be applied to the holdout sample.

`expProfit` requires five input values. First is the name of the `data.table` that contains your holdout sample. Second is the name of the variable in the holdout sample that contains the **score** variables. In this application, this is `predLogitProb`. Third is the name of the variable that contains the actual outcomes that were observed. In this application, this is `churned`. Fourth is the value from the **Cost-Benefit Matrix** of correct predictions that the consumer will take an action. Fifth is the value from the **Cost-Benefit Matrix**

You are not responsible for understanding the coding details in this function, but I am happy to discuss them with you if interested.

```
# Function to calculate the expected profit at each score level
expProfit <- function(DT, score, actual, v11, v10) {
  # prepare copied data set for operations within the function with standardized names
  setnames(DT, c(actual, score), c("actual", "score"))
  DTloc <- copy(DT[, .(actual, score)]) # Copy a new version of the data set to work with in this fun
  setnames(DT, c("actual", "score"), c(actual, score)) # Return original variable names to original d
  DTloc[, origOrder := 1:.N] # Capture the order given for returning the output

  # Sort by score, highest to lowest.
  setorder( DTloc, -score )
  DTloc[, p11 := cumsum(actual==1) /.N] # Probability of true positives at a given score threshold
  DTloc[, p10 := cumsum(actual==0) /.N] # Probability of false positives at a given score threshold
  DTloc[, expProfit := v11*p11 + v10*p10] # Profits

  # Output expected profit estimates
  setorder(DTloc, origOrder)
  return(DTloc$expProfit)
}
```

Apply the `expProfit` function

```
### Base version of the model
# Apply the function to calculate expected profits for each threshold
churn_HO[, expProfitLogitBase := expProfit(
  DT = churn_HO,           # data.table use for the calculation
  score = "predLogitProbBase", # Which variable corresponds to the score
  actual = "churned",      # Which variable corresponds to the actual outcomes
  v11 = 4,                 # Value of a true positive
  v10 = -1)]              # Value of a false positive.

### Version of the model with extra "garbage" variables
# Apply the function to calculate expected profits for each threshold
churn_HO[, expProfitLogitGarb := expProfit(
  DT = churn_HO,           # data.table use for the calculation
  score = "predLogitProbGarb", # Which variable corresponds to the score
  actual = "churned",      # Which variable corresponds to the actual outcomes
  v11 = 4,                 # Value of a true positive
  v10 = -1)]              # Value of a false positive.
```

Estimate using LASSO

We will use the `glmnet` function for LASSO and ridge regressions. `glmnet` does not take data.tables as its input, but instead requires that the X data be converted to a matrix and the y data a vector.

We will use the “cross-validated” version of `glmnet`, which is called `cv.glmnet`. It uses the **cross validation** method we described in class to help discover the optimal value for `lambda`.

```
# Convert the data.table to a matrix and vector to input to glmnet
X = model.matrix( churned ~ . , churn_DT[holdout==0, -..noRegVars] ) # Creates a matrix of the X variables
y = churn_DT[holdout==0, churned] # Creates a vector of the y variables
lasso.fit <- cv.glmnet(X, y, family="binomial", alpha=1.0) # Estimates the model for binary outcomes. alpha=1.0 is LASSO
# (Ridge can be estimated with alpha=0.)
coef(lasso.fit, s = "lambda.min") # Report the estimates for each coefficients estimated
```

317 x 1 sparse Matrix of class "dgCMatrix"

	s1
(Intercept)	-3.813829e+00
(Intercept)	.
genderM	-5.970516e-02
employSelfEmployed	.
employCrafts	.
employProfessional	-2.220383e-01
employClerical	.
employRetired	9.774171e-02
employStudent	.
hhIncome	-1.234730e-06
hhSize	.
homeOwner	.
married	.
avgLatePayment	.
nMonthsCust	.
avgMoBill	2.775950e-02
creditScore	.
X1	.
X2	-4.351818e-02
X3	.
X4	.
X5	.
X6	.
X7	.
X8	.
X9	.
X10	.
X11	.
X12	.
X13	.
X14	.
X15	.
X16	-9.094375e-03
X17	.
X18	.
X19	4.839712e-02
X20	.
X21	.

X22	.
X23	.
X24	8.513802e-02
X25	.
X26	.
X27	.
X28	.
X29	.
X30	.
X31	.
X32	.
X33	.
X34	.
X35	.
X36	.
X37	.
X38	.
X39	.
X40	.
X41	.
X42	.
X43	.
X44	.
X45	.
X46	1.109828e-01
X47	.
X48	.
X49	1.541288e-02
X50	.
X51	.
X52	.
X53	.
X54	.
X55	.
X56	.
X57	.
X58	.
X59	.
X60	.
X61	.
X62	.
X63	.
X64	.
X65	.
X66	.
X67	.
X68	.
X69	.
X70	8.799158e-02
X71	.
X72	.
X73	.
X74	.
X75	.

X76	.
X77	3.533740e-02
X78	.
X79	.
X80	.
X81	.
X82	2.576636e-03
X83	.
X84	.
X85	.
X86	.
X87	.
X88	1.033319e-01
X89	.
X90	.
X91	.
X92	.
X93	.
X94	-4.833706e-03
X95	.
X96	2.333157e-02
X97	.
X98	.
X99	.
X100	.
X101	1.262619e-03
X102	9.272279e-02
X103	.
X104	.
X105	.
X106	.
X107	.
X108	8.462068e-03
X109	.
X110	.
X111	.
X112	.
X113	.
X114	.
X115	.
X116	.
X117	.
X118	.
X119	.
X120	.
X121	.
X122	.
X123	.
X124	.
X125	.
X126	.
X127	.
X128	.
X129	.

X130	.
X131	.
X132	.
X133	.
X134	.
X135	.
X136	-1.014284e-01
X137	.
X138	.
X139	-4.815335e-02
X140	.
X141	.
X142	.
X143	.
X144	.
X145	.
X146	.
X147	.
X148	.
X149	.
X150	.
X151	.
X152	.
X153	.
X154	.
X155	.
X156	.
X157	.
X158	.
X159	8.493787e-02
X160	.
X161	-1.470461e-01
X162	.
X163	.
X164	.
X165	.
X166	.
X167	-7.960225e-02
X168	.
X169	4.623966e-02
X170	5.844784e-05
X171	.
X172	.
X173	.
X174	.
X175	.
X176	.
X177	.
X178	.
X179	.
X180	.
X181	.
X182	.
X183	3.040603e-02

X184	-2.731483e-02
X185	.
X186	.
X187	.
X188	.
X189	.
X190	.
X191	-3.834900e-02
X192	.
X193	.
X194	.
X195	.
X196	1.260179e-01
X197	.
X198	.
X199	.
X200	.
X201	.
X202	.
X203	.
X204	.
X205	.
X206	.
X207	7.062899e-02
X208	.
X209	.
X210	.
X211	.
X212	.
X213	.
X214	.
X215	.
X216	.
X217	.
X218	.
X219	.
X220	.
X221	.
X222	.
X223	.
X224	.
X225	.
X226	.
X227	.
X228	.
X229	.
X230	.
X231	.
X232	.
X233	.
X234	.
X235	.
X236	.
X237	.


X238	-1.033301e-01
X239	-6.660676e-02
X240	-3.570075e-02
X241	.
X242	.
X243	.
X244	.
X245	.
X246	.
X247	-1.890971e-01
X248	.
X249	.
X250	-5.483952e-02
X251	.
X252	.
X253	.
X254	1.758450e-01
X255	-3.400405e-02
X256	.
X257	.
X258	.
X259	.
X260	.
X261	.
X262	-6.914809e-02
X263	-7.893707e-02
X264	.
X265	.
X266	.
X267	.
X268	.
X269	.
X270	.
X271	.
X272	.
X273	.
X274	5.608535e-02
X275	.
X276	.
X277	.
X278	.
X279	.
X280	.
X281	.
X282	.
X283	.
X284	.
X285	.
X286	.
X287	.
X288	.
X289	.
X290	-1.258814e-02
X291	.

```

X292      .
X293      .
X294      .
X295      -1.977191e-01
X296      .
X297      .
X298      .
X299      .
X300      .

```

```
plot(lasso.fit) # Plot the quality of the predictions for different lambda values.
```



Calculate the expected profits based on the LASSO model

```

churn_H0[, predLASSOProb := predict(lasso.fit, chH0mat, s= "lambda.min", type="response")]

# Apply the function to calculate expected profits for each threshold
churn_H0[, expProfitLASSO := expProfit(
  DT = churn_H0,           # data.table use for the calculation
  score = "predLASSOProb", # Which variable corresponds to the score
  actual = "churned",      # Which variable corresponds to the actual outcomes
  v11 = 4,                 # Value of a true positive
  v10 = -1)]              # Value of a false positive.

```

Reformat the data for plotting

ggplot requires that the data be in a specific format for plotting multiple lines. In particular, the variables on the x and y-axes need to each correspond to a single variable name. In this case, x corresponds to **shareTest** and y corresponds to **expProfit**. In order to plot multiple lines on the same plot, we require a third variable that indicates how the different plots are different. In this case, I have called this third variable **method** and

it takes values “LASSO” and “Logit.”

To put the data in the required shape, I have made two new subsets of the original holdout data. The first reports the x and y values for our “Logit” estimates and the second the values for the “LASSO” estimates. These two new data.tables are then stacked one on top of the other. To accomplish this stacking I place each of the two new data.tables as different objects in the list and then call `rbindlist` which “binds” the two data.tables by row.

The resulting data.table is called `profitDT` and can be plot with `ggplot`.

```
# Stack the data from Logit and LASSO profits
chH0logitBase = churn_H0[order(-predLogitProbBase), .(shareTest=(1:.N)/.N, expProfit = expProfitLogitBase)]
chH0logitGarb = churn_H0[order(-predLogitProbGarb), .(shareTest=(1:.N)/.N, expProfit = expProfitLogitGarb)]
chH0lasso = churn_H0[order(-predLASSOProb), .(shareTest=(1:.N)/.N, expProfit = expProfitLASSO, method="LASSO")]

profitDT = rbindlist( list(chH0logitBase, chH0logitGarb, chH0lasso ))

# Plot the two profit curves on one plot
ggplot(data=profitDT, aes(x=shareTest, y= expProfit, color=method)) +
  geom_line() + theme_bw()
```

